

# WashburnPaulHW5

November 13, 2018

## 0.0.1 CSCI E-82 Homework 5 on CNNs

## 0.0.2 Due by 11/13/18 at 11:59pm EST to the Canvas dropbox

## 0.1 This is an individual homework so there should be no collaboration for this homework.

### 0.1.1 Under each problem, we have a place for you to write the answer, or write runnable code that will produce the answer. Show your work.

This is a busy time of year with homework and an exam coming up. We are looking for a successful working result that builds upon the section code and enables you to gain some proficiency with this important and growing field of deep learning.

Depending on your computer, some of the runs may still take a few minutes per epoch. As a result, Problem 4 may take the better part of a day to run, so plan accordingly.

## 0.2 Your Name:

Paul M. Washburn

Some ideas for project:

- trump tweets vs. stock and bond markets
- topic extraction on trump tweets

## 0.3 Dataset

WikiArt is an amazing resource containing centuries of artwork. Since such datasets are wonderful for deep learning, Kaggle has hosted a challenge to characterize the 'fingerprints' of various artists. The Kaggle dataset contains metadata and also a set of images that have been resized so that the shorter dimension is 256 pixels. To make this homework reasonably fast even for those without GPUs, we have further reduced the images to 64 x 64. CNNs and neural networks in general prefer to have consistent sizes. To achieve this, we cut the center 256 pixels from the longer dimension and then shrunk the images by a factor of 4. This isn't a perfect solution since it did cut off a few heads as you will see.

The selected images are for portraits and landscapes. No, we're not talking about the orientation but rather the content of the images. Thanks to help from Rashmi and Dave, we have a small enough data set that should give reasonable results in a timely manner even on just a CPU.

The data were originally divided into a training and a test set. We have further divided the training set into a train and validation set. In this homework you will be using the training set and validation set to train and assess your deep learning models. At the final step, you will see how

well your final training worked on the test set. In each of these directories, there is a truth.txt file that has the image name and whether it is a portrait or landscape scene.

```
In [1]: import warnings
warnings.filterwarnings('ignore')
import tensorflow as tf
import pandas as pd
import numpy as np
import os
import shutil
from matplotlib import pyplot as plt
from tensorflow import keras
from PIL import Image
import glob
from tensorflow.keras.optimizers import *
from tensorflow.keras.layers import *
from tensorflow.keras.models import *
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import backend as K
from tensorflow.keras.callbacks import TensorBoard, EarlyStopping
from time import time
from sklearn.metrics import roc_curve, auc
import seaborn as sns
sns.set(style="whitegrid")
%matplotlib inline

def set_mpl_preferences(ax):
    ax.grid(alpha=.3)
    sns.despine()
    ax.legend(loc='best')
    sns.set(style="whitegrid")

if K.backend()=='tensorflow':
    K.set_image_data_format('channels_last')

print("Tensorflow is installed and is version: ", tf.__version__)
print("Keras is installed and is version: ", tf.keras.__version__)
```

Tensorflow is installed and is version: 1.10.0

Keras is installed and is version: 2.1.6-tf

```
In [2]: class TrainValTensorBoard(TensorBoard):
    def __init__(self, log_dir='./logs/{}'.format(time()), **kwargs):
        # Make the original `TensorBoard` log to a subdirectory 'training'
        training_log_dir = os.path.join(log_dir, 'training')
        super(TrainValTensorBoard, self).__init__(training_log_dir, **kwargs)
```

```

        # Log the validation metrics to a separate subdirectory
        self.val_log_dir = os.path.join(log_dir, 'validation')

    def set_model(self, model):
        # Setup writer for validation metrics
        self.val_writer = tf.summary.FileWriter(self.val_log_dir)
        super(TrainValTensorBoard, self).set_model(model)

    def on_epoch_end(self, epoch, logs=None):
        # Pop the validation logs and handle them separately with
        # `self.val_writer`. Also rename the keys so that they can
        # be plotted on the same figure with the training metrics
        logs = logs or {}
        val_logs = {k.replace('val_', ''): v for k, v in logs.items() if k.startswith('val_')}
        for name, value in val_logs.items():
            summary = tf.Summary()
            summary_value = summary.value.add()
            summary_value.simple_value = value.item()
            summary_value.tag = name
            self.val_writer.add_summary(summary, epoch)
        self.val_writer.flush()

        # Pass the remaining logs to `TensorBoard.on_epoch_end`
        logs = {k: v for k, v in logs.items() if not k.startswith('val_')}
        super(TrainValTensorBoard, self).on_epoch_end(epoch, logs)

    def on_train_end(self, logs=None):
        super(TrainValTensorBoard, self).on_train_end(logs)
        self.val_writer.close()

In [3]: base_dir = 'data/images64'
        train_dir = os.path.join(base_dir, 'train')
        validation_dir = os.path.join(base_dir, 'validation')
        test_dir = os.path.join(base_dir, 'test')

In [4]: train_labels = pd.read_table('data/images64/train/truth.txt', header=None)
        val_labels = pd.read_table('data/images64/validation/truth.txt', header=None)

```

*Do Not Re-run Code Below, Only Run Once*

```

os.mkdir(os.path.join(train_dir, 'landscape'))
os.mkdir(os.path.join(train_dir, 'portrait'))
os.mkdir(os.path.join(validation_dir, 'landscape'))
os.mkdir(os.path.join(validation_dir, 'portrait'))

# move training files
for landscape_pic in train_labels.loc[train_labels[1]=='landscape', 0]:
    current_home = os.path.join(train_dir, landscape_pic)
    new_home = os.path.join(os.path.join(train_dir, 'landscape'), landscape_pic)

```

```

shutil.move(current_home, new_home)

for portrait_pic in train_labels.loc[train_labels[1]=='portrait', 0]:
    current_home = os.path.join(train_dir, portrait_pic)
    new_home = os.path.join(os.path.join(train_dir, 'portrait'), portrait_pic)
    shutil.move(current_home, new_home)

# move validation files
for landscape_pic in val_labels.loc[val_labels[1]=='landscape', 0]:
    current_home = os.path.join(validation_dir, landscape_pic)
    new_home = os.path.join(os.path.join(validation_dir, 'landscape'), landscape_pic)
    shutil.move(current_home, new_home)

for portrait_pic in val_labels.loc[val_labels[1]=='portrait', 0]:
    current_home = os.path.join(validation_dir, portrait_pic)
    new_home = os.path.join(os.path.join(validation_dir, 'portrait'), portrait_pic)
    shutil.move(current_home, new_home)

```

## 0.4 Problem 1 (5 points)


Read in and display the first 5 portraits and the first 5 landscapes. Note, if you are using the OpenCV tools, then the color may be distorted. The `cvtColor()` method using `cv2.COLOR_BGR2RGB` may be useful. However, it is likely easier to use the generator and `plot_strip` example from section.

```

In [5]: train_files = glob.glob('data/images64/train/landscape/*.jpg')
fig, axes = plt.subplots(1, 5, figsize=(17, 5))
for i, f in enumerate(train_files[:5]):
    ax = axes[i]
    img = Image.open(f)
    ax.imshow(np.asarray(img))
    img.close()

plt.show()

```



```

In [6]: train_files = glob.glob('data/images64/train/portrait/*.jpg')
fig, axes = plt.subplots(1, 5, figsize=(17, 5))
for i, f in enumerate(train_files[:5]):


```

```

ax = axes[i]
img = Image.open(f)
ax.imshow(np.asarray(img))
img.close()

plt.show()

```



## 0.5 Problem 2 (25 points)

Construct a baseline CNN classifier using Keras for the training set and assess the validation set performance at each epoch. The goal is to correctly classify portraits from landscapes. Plot the resulting performance on the training and validation set as a function of epoch using the criteria over which you are optimizing. You should run at least 20 epochs for this problem.

```

In [7]: # select batch_size
batch_size = 128

train_datagen = ImageDataGenerator(rescale=1./255,
                                    #featurewise_center=True,
                                    #featurewise_std_normalization=True,
                                    rotation_range=20,
                                    width_shift_range=0.2,
                                    height_shift_range=0.2,
                                    horizontal_flip=True)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(64, 64),
    batch_size=batch_size, # somewhat arbitrarily chosen
    class_mode='binary')

val_datagen = ImageDataGenerator(rescale=1./255,
                                  #featurewise_center=True,
                                  #featurewise_std_normalization=True,
                                  rotation_range=20,
                                  width_shift_range=0.2,
                                  height_shift_range=0.2,
                                  horizontal_flip=True)

validation_generator = val_datagen.flow_from_directory(validation_dir,

```

```
target_size=(64, 64),
batch_size=batch_size,
class_mode='binary')
```

Found 16315 images belonging to 2 classes.

Found 8158 images belonging to 2 classes.

```
In [8]: np.random.seed(777)
```

```
In [9]: K.clear_session()
```

```
pool_size = (3,3)
```

```
early_stopping = EarlyStopping(patience=3)
tensorboard = TensorBoard(log_dir="logs/{}".format(time()))
```

```
model = Sequential(name='cnn')
```

```
# first convolution
```

```
model.add(Conv2D(32, (5, 5), activation='relu',
                input_shape=(64, 64, 3),
                name='conv1',
                padding='same'))
```

```
model.add(AveragePooling2D(pool_size, name='avg_pool1'))
```

```
# second convolution
```

```
model.add(Conv2D(64, (5, 5), activation='relu', name = 'conv2', padding='same'))
model.add(Conv2D(32, (5, 5), activation='relu', name = 'conv3', padding='same'))
model.add(AveragePooling2D(pool_size, name='avg_pool2'))
```

```
model.add(Flatten())
```

```
model.add(Dense(256, activation='relu', name='fc1')) #128
```

```
model.add(Dense(1, activation='sigmoid', name='fc2'))
```

```
sgd = SGD(lr = 0.05, decay=1e-6, momentum=0.9, nesterov=True)
```

```
model.compile(loss='binary_crossentropy',
              #optimizer=optimizers.RMSprop(lr=1e-4),
              #optimizer=sgd,
              optimizer=Adam(lr=1e-4, decay=1e-6),
              metrics=['accuracy'])
```

```
# fit model
```

```
history = model.fit_generator(train_generator,
                             steps_per_epoch=100,
                             epochs=25,
```

```
validation_data=validation_generator,  
validation_steps=50,  
verbose=1,  
callbacks=[tensorboard, early_stopping])
```

```
model.summary()
```

Epoch 1/25

100/100 [=====] - 84s 843ms/step - loss: 0.5473 - acc: 0.7303 - val\_loss: 0.4341

Epoch 2/25

100/100 [=====] - 83s 827ms/step - loss: 0.4341 - acc: 0.8016 - val\_loss: 0.3876

Epoch 3/25

100/100 [=====] - 84s 839ms/step - loss: 0.3876 - acc: 0.8259 - val\_loss: 0.3623

Epoch 4/25

100/100 [=====] - 83s 831ms/step - loss: 0.3623 - acc: 0.8427 - val\_loss: 0.3340

Epoch 5/25

100/100 [=====] - 83s 834ms/step - loss: 0.3340 - acc: 0.8574 - val\_loss: 0.3082

Epoch 6/25

100/100 [=====] - 79s 793ms/step - loss: 0.3306 - acc: 0.8629 - val\_loss: 0.3155

Epoch 7/25

100/100 [=====] - 79s 792ms/step - loss: 0.3082 - acc: 0.8728 - val\_loss: 0.2934

Epoch 8/25

100/100 [=====] - 79s 785ms/step - loss: 0.3155 - acc: 0.8682 - val\_loss: 0.2764

Epoch 9/25

100/100 [=====] - 79s 789ms/step - loss: 0.2934 - acc: 0.8770 - val\_loss: 0.2788

Epoch 10/25

100/100 [=====] - 79s 786ms/step - loss: 0.2974 - acc: 0.8764 - val\_loss: 0.2661

Epoch 11/25

100/100 [=====] - 79s 793ms/step - loss: 0.2822 - acc: 0.8873 - val\_loss: 0.2656

Epoch 12/25

100/100 [=====] - 78s 784ms/step - loss: 0.2735 - acc: 0.8923 - val\_loss: 0.2764

Epoch 13/25

100/100 [=====] - 79s 787ms/step - loss: 0.2764 - acc: 0.8917 - val\_loss: 0.2788

Epoch 14/25

100/100 [=====] - 79s 785ms/step - loss: 0.2788 - acc: 0.8858 - val\_loss: 0.2661

Epoch 15/25

100/100 [=====] - 79s 791ms/step - loss: 0.2661 - acc: 0.8950 - val\_loss: 0.2656

Epoch 16/25

100/100 [=====] - 78s 784ms/step - loss: 0.2656 - acc: 0.8948 - val\_loss: 0.2764

Epoch 17/25

100/100 [=====] - 79s 790ms/step - loss: 0.2685 - acc: 0.8922 - val\_loss: 0.2788

Epoch 18/25

100/100 [=====] - 79s 788ms/step - loss: 0.2628 - acc: 0.8948 - val\_loss: 0.2661

Epoch 19/25

100/100 [=====] - 79s 793ms/step - loss: 0.2591 - acc: 0.8971 - val\_loss: 0.2656

Epoch 20/25

100/100 [=====] - 78s 784ms/step - loss: 0.2516 - acc: 0.8990 - val\_loss: 0.2764

Epoch 21/25

```

100/100 [=====] - 78s 785ms/step - loss: 0.2512 - acc: 0.8991 - val_loss: 0.2512
Epoch 22/25
100/100 [=====] - 79s 787ms/step - loss: 0.2431 - acc: 0.9035 - val_loss: 0.2431
Epoch 23/25
100/100 [=====] - 79s 787ms/step - loss: 0.2454 - acc: 0.9038 - val_loss: 0.2454
Epoch 24/25
100/100 [=====] - 79s 786ms/step - loss: 0.2403 - acc: 0.9058 - val_loss: 0.2403
Epoch 25/25
100/100 [=====] - 79s 785ms/step - loss: 0.2550 - acc: 0.8972 - val_loss: 0.2550

```

Layer (type)	Output Shape	Param #
conv1 (Conv2D)	(None, 64, 64, 32)	896
avg_pool1 (AveragePooling2D)	(None, 21, 21, 32)	0
conv2 (Conv2D)	(None, 21, 21, 64)	51264
conv3 (Conv2D)	(None, 21, 21, 32)	51232
avg_pool2 (AveragePooling2D)	(None, 7, 7, 32)	0
flatten (Flatten)	(None, 1568)	0
fc1 (Dense)	(None, 256)	401664
fc2 (Dense)	(None, 1)	257
Total params: 505,313		
Trainable params: 505,313		
Non-trainable params: 0		

```
In [10]: model.save_weights('models/attempt_0.h5')
```

To start tensorboard:

```
Pauls-MacBook-Pro:HW5 pmw$ tensorboard --logdir=logs
```

## 0.6 Problem 3 (5 points)

From the pattern of training and validation curves, describe what is good/bad and what you plan to do next to improve the result.

## 0.7 Problem 4 (45 points)

This step is where we want you to do most of your personal learning. Your goal is to improve the network using a combination of architecture choices, parameter tuning, and experimenting



with different optimizers/dropout/regularization/etc. Treat each of these as separate optimization/exploration steps for now. We would like to see 3 separate steps that cover different areas. The format of the 3 steps should be as follows: \* State the hypothesis/strategy for how you will improve/explore a particular aspect. \* Describe what types of tests you are running and why (i.e. what range of parameters are you choosing and why) \* Include the code and results \* State your interpretation of the results

We're not looking for research in deep learning, but we want you to gain some hands-on experience working with Keras and figuring out what works. A good example may be comparing strategies to overcome overfitting, or comparing a few different CNN architectures in terms of performance and speed, or comparing data augmentation types and results.

### 0.7.1 Hypothesis 1: Adding another full convolution layer

```
In [11]: np.random.seed(777)
```

```
K.clear_session()

pool_size = (3, 3)

early_stopping = EarlyStopping(patience=3)
tensorboard = TensorBoard(log_dir="logs/{}".format(time()))

model = Sequential(name='cnn')

# first convolution
model.add(Conv2D(32, (5, 5), activation='relu',
                input_shape=(64, 64, 3),
                name='conv1',
                padding='same'))
model.add(AveragePooling2D(pool_size, name='avg_pool1'))

# second convolution
model.add(Conv2D(64, (5, 5), activation='relu', name = 'conv2', padding='same'))
model.add(Conv2D(32, (5, 5), activation='relu', name = 'conv3', padding='same'))
model.add(AveragePooling2D(pool_size, name='avg_pool2'))

# third - new - convolution
model.add(Conv2D(64, (5, 5), activation='relu', name = 'conv4', padding='same'))
model.add(Conv2D(32, (5, 5), activation='relu', name = 'conv5', padding='same'))
model.add(AveragePooling2D(pool_size, name='avg_pool3'))

model.add(Flatten())
model.add(Dense(256, activation='relu', name='fc1'))
model.add(Dense(1, activation='sigmoid', name='fc2'))

sgd = SGD(lr = 0.05, decay=1e-6, momentum=0.9, nesterov=True)

model.compile(loss='binary_crossentropy',
```

```

#optimizer=optimizers.RMSprop(lr=1e-4),
#optimizer=sgd,
optimizer=Adam(lr=1e-4, decay=1e-6),
metrics=['accuracy'])

```

```

# fit model

```

```

history = model.fit_generator(train_generator,
                              steps_per_epoch=100,
                              epochs=25,
                              validation_data=validation_generator,
                              validation_steps=50,
                              verbose=1,
                              callbacks=[tensorboard, early_stopping])

```

```

model.summary()

```

```

Epoch 1/25
100/100 [=====] - 85s 855ms/step - loss: 0.5893 - acc: 0.7049 - val_loss: 0.4429
Epoch 2/25
100/100 [=====] - 86s 859ms/step - loss: 0.4429 - acc: 0.7985 - val_loss: 0.3886
Epoch 3/25
100/100 [=====] - 82s 823ms/step - loss: 0.3886 - acc: 0.8257 - val_loss: 0.3762
Epoch 4/25
100/100 [=====] - 84s 841ms/step - loss: 0.3762 - acc: 0.8370 - val_loss: 0.3518
Epoch 5/25
100/100 [=====] - 83s 834ms/step - loss: 0.3518 - acc: 0.8489 - val_loss: 0.3418
Epoch 6/25
100/100 [=====] - 83s 827ms/step - loss: 0.3418 - acc: 0.8550 - val_loss: 0.3257
Epoch 7/25
100/100 [=====] - 83s 828ms/step - loss: 0.3257 - acc: 0.8648 - val_loss: 0.3170
Epoch 8/25
100/100 [=====] - 84s 840ms/step - loss: 0.3170 - acc: 0.8723 - val_loss: 0.3145
Epoch 9/25
100/100 [=====] - 84s 837ms/step - loss: 0.3145 - acc: 0.8670 - val_loss: 0.2963
Epoch 10/25
100/100 [=====] - 83s 834ms/step - loss: 0.2963 - acc: 0.8798 - val_loss: 0.2983
Epoch 11/25
100/100 [=====] - 83s 833ms/step - loss: 0.2983 - acc: 0.8791 - val_loss: 0.2995
Epoch 12/25
100/100 [=====] - 84s 840ms/step - loss: 0.2995 - acc: 0.8748 - val_loss: 0.2854
Epoch 13/25
100/100 [=====] - 84s 836ms/step - loss: 0.2854 - acc: 0.8838 - val_loss: 0.2854

```

Layer (type)	Output Shape	Param #
conv1 (Conv2D)	(None, 64, 64, 32)	896

avg_pool1 (AveragePooling2D)	(None, 21, 21, 32)	0
-----		
conv2 (Conv2D)	(None, 21, 21, 64)	51264
-----		
conv3 (Conv2D)	(None, 21, 21, 32)	51232
-----		
avg_pool2 (AveragePooling2D)	(None, 7, 7, 32)	0
-----		
conv4 (Conv2D)	(None, 7, 7, 64)	51264
-----		
conv5 (Conv2D)	(None, 7, 7, 32)	51232
-----		
avg_pool3 (AveragePooling2D)	(None, 2, 2, 32)	0
-----		
flatten (Flatten)	(None, 128)	0
-----		
fc1 (Dense)	(None, 256)	33024
-----		
fc2 (Dense)	(None, 1)	257
=====		
Total params: 239,169		
Trainable params: 239,169		
Non-trainable params: 0		
-----		

```
In [12]: model.save_weights('models/attempt_1.h5')
```

## 0.7.2 Hypothesis 2: Using MaxPooling2D will improve the accuracy

```
In [13]: K.clear_session()
```

```
kernel_size = (3, 3)

early_stopping = EarlyStopping(patience=3)
tensorboard = TensorBoard(log_dir="logs/{}".format(time()))

model = Sequential(name='cnn')

# first convolution
model.add(Conv2D(32, kernel_size, activation='relu',
                 input_shape=(64, 64, 3),
                 name='conv1',
                 padding='same'))
model.add(MaxPooling2D(kernel_size, name='max_pool1'))

# second convolution
model.add(Conv2D(64, (5, 5), activation='relu', name = 'conv2', padding='same'))
```

```

model.add(Conv2D(32, (5, 5), activation='relu', name = 'conv3', padding='same'))
model.add(MaxPooling2D(kernel_size, name='max_pool2'))

# third - new - convolution
model.add(Conv2D(64, (5, 5), activation='relu', name = 'conv4', padding='same'))
model.add(Conv2D(32, (5, 5), activation='relu', name = 'conv5', padding='same'))
model.add(AveragePooling2D(kernel_size, name='max_pool3'))

model.add(Flatten())
model.add(Dense(256, activation='relu', name='fc1'))
model.add(Dense(1, activation='sigmoid', name='fc2'))

sgd = SGD(lr = 0.05, decay=1e-6, momentum=0.9, nesterov=True)

model.compile(loss='binary_crossentropy',
              #optimizer=optimizers.RMSprop(lr=1e-4),
              optimizer=Adam(lr=1e-4, decay=1e-6),#sgd,
              metrics=['accuracy'])

# fit model
history = model.fit_generator(train_generator,
                             steps_per_epoch=100,
                             epochs=25,
                             validation_data=validation_generator,
                             validation_steps=50,
                             verbose=1,
                             callbacks=[tensorboard, early_stopping])

model.summary()

```

```

Epoch 1/25
100/100 [=====] - 86s 862ms/step - loss: 0.5599 - acc: 0.7376 - val_loss: 0.4111
Epoch 2/25
100/100 [=====] - 84s 843ms/step - loss: 0.3876 - acc: 0.8279 - val_loss: 0.3111
Epoch 3/25
100/100 [=====] - 85s 855ms/step - loss: 0.3564 - acc: 0.8450 - val_loss: 0.2889
Epoch 4/25
100/100 [=====] - 84s 841ms/step - loss: 0.3479 - acc: 0.8503 - val_loss: 0.2889
Epoch 5/25
100/100 [=====] - 84s 842ms/step - loss: 0.3194 - acc: 0.8694 - val_loss: 0.2889
Epoch 6/25
100/100 [=====] - 85s 851ms/step - loss: 0.3107 - acc: 0.8724 - val_loss: 0.2889
Epoch 7/25
100/100 [=====] - 83s 835ms/step - loss: 0.2884 - acc: 0.8840 - val_loss: 0.2889
Epoch 8/25
100/100 [=====] - 84s 845ms/step - loss: 0.3034 - acc: 0.8766 - val_loss: 0.2889
Epoch 9/25

```

```

100/100 [=====] - 84s 842ms/step - loss: 0.2915 - acc: 0.8807 - val_loss: 0.2915
Epoch 10/25
100/100 [=====] - 85s 849ms/step - loss: 0.2691 - acc: 0.8953 - val_loss: 0.2691
Epoch 11/25
100/100 [=====] - 84s 836ms/step - loss: 0.2639 - acc: 0.8955 - val_loss: 0.2639
Epoch 12/25
100/100 [=====] - 84s 841ms/step - loss: 0.2672 - acc: 0.8912 - val_loss: 0.2672
Epoch 13/25
100/100 [=====] - 84s 844ms/step - loss: 0.2522 - acc: 0.8990 - val_loss: 0.2522
Epoch 14/25
100/100 [=====] - 83s 831ms/step - loss: 0.2526 - acc: 0.8985 - val_loss: 0.2526
Epoch 15/25
100/100 [=====] - 85s 845ms/step - loss: 0.2456 - acc: 0.8986 - val_loss: 0.2456
Epoch 16/25
100/100 [=====] - 83s 830ms/step - loss: 0.2308 - acc: 0.9061 - val_loss: 0.2308
Epoch 17/25
100/100 [=====] - 86s 864ms/step - loss: 0.2310 - acc: 0.9067 - val_loss: 0.2310
Epoch 18/25
100/100 [=====] - 84s 837ms/step - loss: 0.2297 - acc: 0.9082 - val_loss: 0.2297
Epoch 19/25
100/100 [=====] - 84s 840ms/step - loss: 0.2270 - acc: 0.9103 - val_loss: 0.2270
Epoch 20/25
100/100 [=====] - 85s 846ms/step - loss: 0.2312 - acc: 0.9065 - val_loss: 0.2312
Epoch 21/25
100/100 [=====] - 83s 835ms/step - loss: 0.2202 - acc: 0.9100 - val_loss: 0.2202
Epoch 22/25
100/100 [=====] - 84s 842ms/step - loss: 0.2200 - acc: 0.9080 - val_loss: 0.2200
Epoch 23/25
100/100 [=====] - 84s 841ms/step - loss: 0.2132 - acc: 0.9129 - val_loss: 0.2132
Epoch 24/25
100/100 [=====] - 84s 844ms/step - loss: 0.2092 - acc: 0.9145 - val_loss: 0.2092
Epoch 25/25
100/100 [=====] - 84s 837ms/step - loss: 0.2173 - acc: 0.9109 - val_loss: 0.2173

```

Layer (type)	Output Shape	Param #
conv1 (Conv2D)	(None, 64, 64, 32)	896
max_pool1 (MaxPooling2D)	(None, 21, 21, 32)	0
conv2 (Conv2D)	(None, 21, 21, 64)	51264
conv3 (Conv2D)	(None, 21, 21, 32)	51232
max_pool2 (MaxPooling2D)	(None, 7, 7, 32)	0
conv4 (Conv2D)	(None, 7, 7, 64)	51264

conv5 (Conv2D)	(None, 7, 7, 32)	51232
-----		
max_pool3 (AveragePooling2D)	(None, 2, 2, 32)	0
-----		
flatten (Flatten)	(None, 128)	0
-----		
fc1 (Dense)	(None, 256)	33024
-----		
fc2 (Dense)	(None, 1)	257
=====		
Total params: 239,169		
Trainable params: 239,169		
Non-trainable params: 0		
-----		

```
In [14]: model.save_weights('models/attempt_2.h5')
```

### 0.7.3 Hypothesis 3: kernel\_size = (7, 7) will improve the model

```
In [15]: K.clear_session()
```

```
kernel_size = (7, 7)
pool_size = (2, 2)

early_stopping = EarlyStopping(patience=2)
tensorboard = TensorBoard(log_dir="logs/{}".format(time()))

model = Sequential(name='cnn')

# first convolution
model.add(Conv2D(32, kernel_size, activation='relu',
                 input_shape=(64, 64, 3),
                 name='conv1',
                 padding='same'))
model.add(MaxPooling2D(pool_size, name='max_pool1'))

# second convolution
model.add(Conv2D(64, kernel_size, activation='relu', name = 'conv2', padding='same'))
model.add(Conv2D(32, kernel_size, activation='relu', name = 'conv3', padding='same'))
model.add(MaxPooling2D(pool_size, name='max_pool2'))

# third - new - convolution
model.add(Conv2D(64, kernel_size, activation='relu', name = 'conv4', padding='same'))
model.add(Conv2D(32, kernel_size, activation='relu', name = 'conv5', padding='same'))
model.add(AveragePooling2D(pool_size, name='max_pool3'))

model.add(Flatten())
```

```

model.add(Dense(256, activation='relu', name='fc1'))
model.add(Dense(1, activation='sigmoid', name='fc2'))

sgd = SGD(lr = 0.05, decay=1e-6, momentum=0.9, nesterov=True)

model.compile(loss='binary_crossentropy',
              #optimizer=RMSprop(lr=1e-4),
              optimizer=Adam(lr=1e-4, decay=1e-6), #sgd,
              metrics=['accuracy'])

# fit model
history = model.fit_generator(train_generator,
                             steps_per_epoch=100,
                             epochs=25,
                             validation_data=validation_generator,
                             validation_steps=50,
                             verbose=1,
                             callbacks=[tensorboard, early_stopping])

model.summary()

```

Using TensorFlow backend.

```

Epoch 1/25
100/100 [=====] - 310s 3s/step - loss: 0.5128 - acc: 0.7495 - val_loss: 0.4000
Epoch 2/25
100/100 [=====] - 310s 3s/step - loss: 0.3852 - acc: 0.8261 - val_loss: 0.3000
Epoch 3/25
100/100 [=====] - 311s 3s/step - loss: 0.3441 - acc: 0.8536 - val_loss: 0.2500
Epoch 4/25
100/100 [=====] - 326s 3s/step - loss: 0.3153 - acc: 0.8680 - val_loss: 0.2000
Epoch 5/25
100/100 [=====] - 327s 3s/step - loss: 0.3027 - acc: 0.8728 - val_loss: 0.1500
Epoch 6/25
100/100 [=====] - 339s 3s/step - loss: 0.2771 - acc: 0.8888 - val_loss: 0.1000
Epoch 7/25
100/100 [=====] - 326s 3s/step - loss: 0.2657 - acc: 0.8916 - val_loss: 0.0500
Epoch 8/25
100/100 [=====] - 319s 3s/step - loss: 0.2668 - acc: 0.8903 - val_loss: 0.0000
Epoch 9/25
100/100 [=====] - 321s 3s/step - loss: 0.2407 - acc: 0.9032 - val_loss: 0.0000
Epoch 10/25
100/100 [=====] - 308s 3s/step - loss: 0.2372 - acc: 0.9057 - val_loss: 0.0000
Epoch 11/25
100/100 [=====] - 310s 3s/step - loss: 0.2284 - acc: 0.9068 - val_loss: 0.0000
Epoch 12/25

```

```

100/100 [=====] - 318s 3s/step - loss: 0.2260 - acc: 0.9061 - val_loss: 0.2260
Epoch 13/25
100/100 [=====] - 308s 3s/step - loss: 0.2190 - acc: 0.9109 - val_loss: 0.2190
Epoch 14/25
100/100 [=====] - 309s 3s/step - loss: 0.2118 - acc: 0.9134 - val_loss: 0.2118
Epoch 15/25
100/100 [=====] - 310s 3s/step - loss: 0.2148 - acc: 0.9132 - val_loss: 0.2148
Epoch 16/25
100/100 [=====] - 308s 3s/step - loss: 0.2122 - acc: 0.9134 - val_loss: 0.2122
Epoch 17/25
100/100 [=====] - 312s 3s/step - loss: 0.2063 - acc: 0.9189 - val_loss: 0.2063
Epoch 18/25
100/100 [=====] - 321s 3s/step - loss: 0.2039 - acc: 0.9162 - val_loss: 0.2039
Epoch 19/25
100/100 [=====] - 327s 3s/step - loss: 0.2015 - acc: 0.9206 - val_loss: 0.2015
Epoch 20/25
100/100 [=====] - 317s 3s/step - loss: 0.2031 - acc: 0.9155 - val_loss: 0.2031
Epoch 21/25
100/100 [=====] - 310s 3s/step - loss: 0.1967 - acc: 0.9213 - val_loss: 0.1967

```

Layer (type)	Output Shape	Param #
conv1 (Conv2D)	(None, 64, 64, 32)	4736
max_pool1 (MaxPooling2D)	(None, 32, 32, 32)	0
conv2 (Conv2D)	(None, 32, 32, 64)	100416
conv3 (Conv2D)	(None, 32, 32, 32)	100384
max_pool2 (MaxPooling2D)	(None, 16, 16, 32)	0
conv4 (Conv2D)	(None, 16, 16, 64)	100416
conv5 (Conv2D)	(None, 16, 16, 32)	100384
max_pool3 (AveragePooling2D)	(None, 8, 8, 32)	0
flatten (Flatten)	(None, 2048)	0
fc1 (Dense)	(None, 256)	524544
fc2 (Dense)	(None, 1)	257

```

Total params: 931,137
Trainable params: 931,137
Non-trainable params: 0

```



```
In [16]: model.save_weights('models/attempt_3.h5')
```

## 0.8 Problem 5 (10 points)

Assess your best model on the test data. Plot the corresponding ROC curve from the results (since we've provided the truth). This was not directly covered in section, but will require a prediction using images in the same format as the training. We suggest referring to the Keras API else use a Google to search to find how to make predictions.

```
In [17]: test_labels = pd.read_table('data/images64/test/truth.txt', header=None)
        test_labels.head()
```

```
Out[17]:
```

	0	1
0	3.jpg	portrait
1	34.jpg	portrait
2	68.jpg	portrait
3	87.jpg	landscape
4	92.jpg	portrait

```
In [18]: test_labels[1].value_counts() / test_labels.shape[0]
```

```
Out[18]:
```

portrait	0.531373
landscape	0.468627

Name: 1, dtype: float64

```
# run this only once
```

```
os.mkdir(os.path.join(test_dir, 'landscape'))
os.mkdir(os.path.join(test_dir, 'portrait'))
```

```
# move test files
```

```
for landscape_pic in test_labels.loc[test_labels[1]=='landscape', 0]:
    current_home = os.path.join(test_dir, landscape_pic)
    new_home = os.path.join(os.path.join(test_dir, 'landscape'), landscape_pic)
    shutil.move(current_home, new_home)
```

```
for portrait_pic in test_labels.loc[test_labels[1]=='portrait', 0]:
    current_home = os.path.join(test_dir, portrait_pic)
    new_home = os.path.join(os.path.join(test_dir, 'portrait'), portrait_pic)
    shutil.move(current_home, new_home)
```

```
In [67]: batch_size = 1
```

```
test_datagen = ImageDataGenerator(rescale=1./255)
test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(64, 64),
    batch_size=128, # somewhat arbitrarily chosen)
```

```
class_mode='binary',  
shuffle=False)
```

```
phat_test = model.predict_generator(test_generator)  
yhat_test = np.array([p > .5 for p in phat_test])
```

Found 7379 images belonging to 2 classes.

```
In [68]: y_test = test_generator.classes  
y_test
```


```
Out[68]: array([0, 0, 0, ..., 1, 1, 1], dtype=int32)
```

```
In [69]: # Compute ROC curve and Area Under the Curve  
fpr, tpr, thresholds = roc_curve(y_test, phat_test)  
roc_auc = auc(fpr, tpr)
```

```
In [70]: print(''  
Area Under Curve (AUC) = %.4f  
''%roc_auc)
```

Area Under Curve (AUC) = 0.9823

```
In [71]: fig, ax = plt.subplots(figsize=(12, 7))  
ax.plot(fpr, tpr, linestyle=':', label='ROC Curve')  
ax.plot(np.arange(0, 1, .01), np.arange(0, 1, .01), label='Random Chance',  
        linestyle='-.', alpha=.5, color='black')  
set_mpl_preferences(ax)  
ax.set_title('ROC Curve for Best Model', size=18)  
ax.set_xlabel('False Positive Rate')  
ax.set_ylabel('True Positive Rate')  
plt.show()
```



```
In [74]: from sklearn.metrics import accuracy_score, f1_score
```

```
acc = accuracy_score(y_test, yhat_test)
f1 = f1_score(y_test, yhat_test)
```

```
print('''
Accuracy Score = %.4f
F1 Score = %.4f
''')
```

```
Accuracy Score = 0.9331
```

```
F1 Score = 0.9365
```

The model looks pretty good! High testing accuracy score and F1 score indicate a robust model. Of course, only time will tell if this is actually the case.

## 0.9 Problem 6 (5 points)

Display the 5 images [worst] misclassified images for each class. Worst is in brackets since certain architectures may only make a binary decision rather than a score. In that case, plot 5 of each.

```
In [78]: test_labels['proba'] = phat_test
test_labels['yhat'] = [int(y) for y in yhat_test]
```

```
test_labels['y'] = y_test
test_labels.head()
```

```
Out [78]:
```

	0	1	proba	yhat	y
0	3.jpg	portrait	0.001469	0	0
1	34.jpg	portrait	0.001025	0	0
2	68.jpg	portrait	0.000299	0	0
3	87.jpg	landscape	0.014885	0	0
4	92.jpg	portrait	0.000164	0	0

```
In [83]: is_wrong = test_labels['yhat'] != test_labels['y']
is_landscape = test_labels[1] == 'landscape'
is_portrait = test_labels[1] == 'portrait'
```

```
Out [83]:
```

	0	1	proba	yhat	y
4655	65550.jpg	landscape	0.000490	0	1
6840	95988.jpg	landscape	0.002054	0	1
5593	79396.jpg	landscape	0.003342	0	1
4675	66023.jpg	landscape	0.006909	0	1
4060	57302.jpg	landscape	0.008769	0	1


```
In [90]: misclassified_landscapes = test_labels.loc[is_wrong & is_landscape]
misclassified_landscapes.sort_values('proba', ascending=True, inplace=True)
misclassified_landscapes.head()
```

```
Out [90]:
```

	0	1	proba	yhat	y
4655	65550.jpg	landscape	0.000490	0	1
6840	95988.jpg	landscape	0.002054	0	1
5593	79396.jpg	landscape	0.003342	0	1
4675	66023.jpg	landscape	0.006909	0	1
4060	57302.jpg	landscape	0.008769	0	1

```
In [89]: files = ['data/images64/test/landscape/'+str(f) for f in misclassified_landscapes[0].index]
fig, axes = plt.subplots(1, 5, figsize=(17, 5))
for i, f in enumerate(files):
    ax = axes[i]
    img = Image.open(f)
    ax.imshow(np.asarray(img))
    img.close()
```

```
plt.show()
```




```
In [92]: misclassified_portraits = test_labels.loc[is_wrong & is_portrait]
misclassified_portraits.sort_values('proba', ascending=False, inplace=True)
misclassified_portraits.head()
```

```
Out[92]:
```

	0	1	proba	yhat	y
431	5650.jpg	portrait	0.996865	1	0
1862	26252.jpg	portrait	0.985408	1	0
1537	21756.jpg	portrait	0.978881	1	0
2159	30115.jpg	portrait	0.976010	1	0
1759	24966.jpg	portrait	0.974267	1	0

```
In [94]: files = ['data/images64/test/portrait/'+str(f) for f in misclassified_portraits[0].head().index]
fig, axes = plt.subplots(1, 5, figsize=(17, 5))
for i, f in enumerate(files):
    ax = axes[i]
    img = Image.open(f)
    ax.imshow(np.asarray(img))
    img.close()
```

```
plt.show()
```



## 0.10 Problem 7 (2 points)

How many hours did this homework take you? The answer to this question will not affect your grade.

```
In [80]: print(''
About 12 hours.
'')
```

About 12 hours.

### 0.11 Last step (3 points)

Save this notebook as LastnameFirstnameHW5.ipynb such as PriceDavid.ipynb. Create a pdf of this notebook named similarly. Submit both the python notebook and the pdf version to the Canvas dropbox. We require both versions.

```
In [81]: print(''  
         OK!  
         '' )
```

OK!